

Deep Learning: Image Classification with CNN

A wide-angle photograph of a city street at dusk or dawn. The street is multi-laned and has several cars. Tall buildings line both sides of the street, with some lights on. The sky is dark. Overlaid on the image are numerous white bounding boxes of varying sizes, each containing a red number. These numbers represent classification scores or confidence levels for different objects in the scene. For example, a car in the foreground has a score of 100, a car further down the road has a score of 10, and a car on the right side of the road has a score of 120. Other numbers visible include 45, 50, and 100.

Collaborators:

- Paola Rivera
- Javier Dastas

Data Processing

```
dataset = '/kaggle/input/animals10'
```

```
Total number of images: 26180  
DataSet Shapes: torch.Size([128, 128, 3])
```

```
Training set size: 22253  
Validation set size: 3927
```

```
Class names: ['cane', 'cavallo', 'elefante', 'farfalla', 'gallina', 'gatto', 'mucca', 'pecora', 'ragno', 'scoiattolo']
```

```
Classes in training set: {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}  
Classes in validation set: {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

< > data	
Name	
>	butterfly
>	cat
>	chicken
>	cow
>	dog
>	elephant
>	horse
>	sheep
>	spider
>	squirrel

```
label  
cane                4863  
ragno               4821  
gallina             3098  
cavallo             2623  
farfalla            2112  
mucca               1866  
scoiattolo          1862  
pecora              1820  
gatto               1668  
elefante            1446  
.ipynb_checkpoints  1  
Name: count, dtype: int64
```

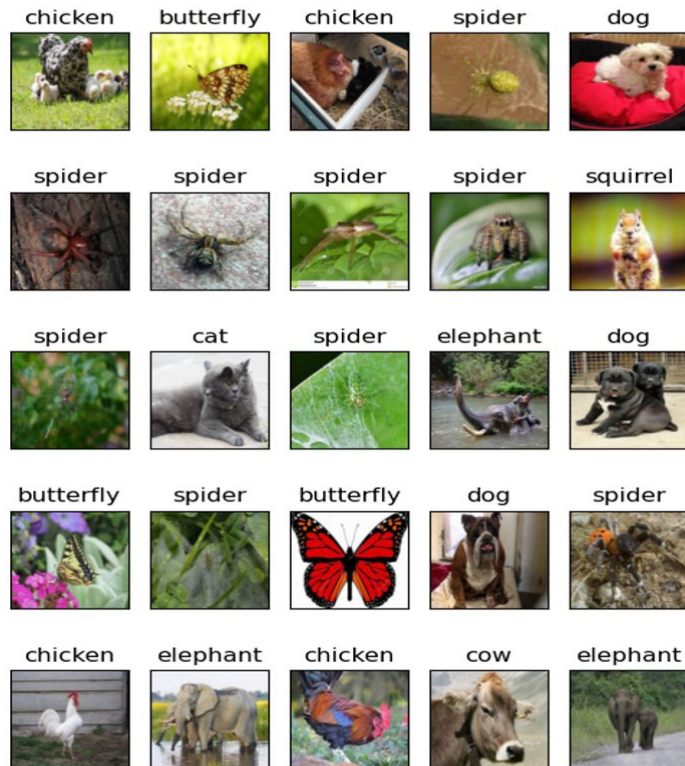
```
Loaded 26180 images with shape (128, 128, 3)
```

Data Processing

```
# Visualize some images
batch_images, batch_labels = next(train_generator)
plt.figure(figsize=(5, 7))
for i in range(25):
    plt.subplot(5, 5, i + 1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(batch_images[i])
    plt.xlabel(list(train_generator.class_indices.keys())[np.argmax(batch_labels[i])])
plt.show()
```

label	
cane	4863
ragno	4821
gallina	3098
cavallo	2623
farfalla	2112
mucca	1866
scoiattolo	1862
pecora	1820
gatto	1668
elefante	1446
.ipynb_checkpoints	1
Name: count, dtype: int64	

Loaded 26180 images with shape (128, 128, 3)



Model Architecture

Input Layer: Accepts images of size 128x128x3 (RGB images).

Convolutional Blocks:

- **3 blocks of convolutional layers** extract spatial features.
- Each block has **Conv2D** layers with **ReLU** activation and **BatchNormalization** to stabilize learning.
- **MaxPooling** reduces spatial dimensions, focusing on key features.

Global Average Pooling: Reduces feature maps to a single vector, minimizing overfitting.

Fully Connected Layers:

- **Dense layer** (128 neurons) learns complex patterns.
- **Dropout layers** (0.1 and 0.5 rates) prevent overfitting.
- **Final Dense layer** (10 neurons with softmax) predicts probabilities for 10 classes.

Regularization: **BatchNormalization** and **Dropout** improve training stability and generalization.

Model Training

- **Steps Per Epoch:**
 - Calculated based on the number of batches (12) in the training and validation datasets to ensure all data is processed in each epoch.
- **Model Compilation:**
 - Optimizer: Adam (adaptive learning rate).
 - Loss Function: Sparse categorical crossentropy for multi-class classification.
 - Metrics: Accuracy to evaluate model performance.
- **Early Stopping:**
 - Monitors validation loss to prevent overfitting.
 - Stops training if validation loss doesn't improve for 5 consecutive epochs.
 - Restores the best model weights from the training process.
- **Training:**
 - Runs for 12 epochs, with the training and validation datasets passed via data loaders.
 - History logs metrics like loss and accuracy for both training and validation sets after each epoch.
- **Results:**
 - Training accuracy improves steadily, reaching 78.6%, while validation accuracy peaks at 75.1%.
 - Validation loss decreases initially but fluctuates, indicating potential overfitting.

Model Evaluation

Validation Performance:

- Loss and accuracy are evaluated on the validation dataset, yielding:
 - Validation Loss: *value from evaluation output.*
 - Validation Accuracy: *value from evaluation output.*

Predictions and Ground Truth:

- Predictions are generated using the validation dataset.
- True Labels: Extracted directly from the validation loader.
- Validated that the lengths of true and predicted labels match.

Classification Metrics:

- Classification Report: Summarizes precision, recall, F1-score, and support for each class using class labels.
- Confusion Matrix: Highlights misclassifications and provides a class-wise error breakdown.

Training History Visualization:

- Loss Plot: Tracks training and validation loss across epochs to identify overfitting or underfitting.
- Accuracy Plot: Shows training and validation accuracy trends, highlighting performance improvements.

Model Evaluation

Len true_classes: 3927

Len predicted_classes: 3927

Classification Report:



	precision	recall	f1-score	support
cane	0.79	0.73	0.76	701
cavallo	0.75	0.67	0.71	422
elefante	0.84	0.70	0.77	220
farfalla	0.68	0.87	0.76	296
gallina	0.85	0.78	0.82	486
gatto	0.73	0.58	0.65	226
mucca	0.72	0.58	0.65	289
pecora	0.61	0.75	0.67	277
ragno	0.74	0.94	0.83	720
scoiattolo	0.80	0.62	0.70	290
accuracy			0.75	3927
macro avg	0.75	0.72	0.73	3927
weighted avg	0.76	0.75	0.75	3927

Confusion Matrix:

```
[[512 31 4 21 16 21 14 42 29 11]
 [ 29 284 12 10 11 1 31 18 24 2]
 [ 9 15 155 3 2 3 4 16 10 3]
 [ 2 0 0 257 2 1 0 0 33 1]
 [ 21 3 2 26 379 3 0 8 36 8]
 [ 42 1 1 12 2 132 2 3 27 4]
 [ 16 33 2 4 7 2 169 36 13 7]
 [ 4 8 6 2 9 5 13 207 16 7]
 [ 6 1 2 29 2 4 0 0 675 1]
 [ 11 2 0 14 14 8 1 9 52 179]]
```

Validation Loss: 0.8091505169868469

Validation Accuracy: 0.7509549260139465



Transfer Learning (Pre-Trained Models)

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
VGG16	528	71.3%	90.1%	138.4M	16	69.5	4.2
VGG19	549	71.3%	90.0%	143.7M	19	84.8	4.4
ResNet50	98	74.9%	92.1%	25.6M	107	58.2	4.6
ResNet50V2	98	76.0%	93.0%	25.6M	103	45.6	4.4
ResNet101	171	76.4%	92.8%	44.7M	209	89.6	5.2
ResNet101V2	171	77.2%	93.8%	44.7M	205	72.7	5.4
ResNet152	232	76.6%	93.1%	60.4M	311	127.4	6.5
ResNet152V2	232	78.0%	94.2%	60.4M	307	107.5	6.6
InceptionV3	92	77.9%	93.7%	23.9M	189	42.2	6.9
InceptionResNetV2	215	80.3%	95.3%	55.9M	449	130.2	10.0

Transfer Learning (Freezing)

Pre-Trained Model:

- Used EfficientNetB0 (pre-trained on ImageNet) as the base model.
- The base model's weights are frozen, retaining its learned features.

Custom Classification Head:

- Added layers on top of the base model to adapt it to the current task:
 - Global Average Pooling: Reduces the spatial dimensions of feature maps.
 - Dense Layer: 128 neurons with ReLU activation for feature learning.
 - Dropout (0.5): Prevents overfitting.
 - Output Layer: 10 neurons with softmax activation for multi-class classification.

Model Compilation:

- Optimizer: Adam with a learning rate of 0.001.
- Loss Function: Sparse categorical crossentropy for classification.
- Metric: Accuracy to evaluate model performance.

Transfer Learning - Data Set Preparation

- **Data Augmentation:**

- Applied to the training set to improve model generalization and handle overfitting:
 - Normalization: Scales pixel values to $[0, 1]$.
 - Transformations: Includes random rotations, shifts, zoom, shear, and horizontal flips.
- Validation data is normalized but not augmented to evaluate the model fairly.

- **Data Splitting:**

- Used ImageDataGenerator to split the dataset:
 - 80% Training: For model learning.
 - 20% Validation: For performance evaluation.

- **Class Imbalance Handling:**

- Class Weights: Computed using `compute_class_weight` to address imbalanced class distribution.
- Ensures the model gives equal importance to underrepresented classes during training.
- Example: Class weights dynamically adjust the loss contribution for each class.

Transfer Learning (Unfreezing)

Fine-Tuning:

- Fine-tuning was performed twice, with improved results in the second iteration.

Unfreezing the Base Model:

- The EfficientNetB0 base model was unfrozen, allowing its pre-trained layers to be updated during training.

Recompilation:

- Lower Learning Rate: Reduced to 0.0001 to avoid large weight updates and maintain the pre-trained knowledge.
- Retained Sparse Categorical Crossentropy as the loss function and Accuracy as the metric.

Training Process:

- Epochs: Fine-tuned for 5 additional epochs.
- Used class weights to address class imbalances in the dataset.
- Training and validation steps were calculated based on the dataset size.

Key Outcome:

- Fine-tuning enhanced the model's performance, leveraging both pre-trained features and task-specific updates.

Model Deployment

164/164 ————— **12s** 71ms/step – accuracy: 0.2661 – loss: 3.4453

Validation Loss: 3.4260590076446533

Validation Accuracy: 0.2696865499019623

164/164 ————— **13s** 74ms/step

Len true_classes: 5232

Len predicted_classes: 5232


Classification Report:

	precision	recall	f1-score	support
cane	0.12	0.01	0.03	972
cavallo	0.05	0.00	0.00	524
elefante	0.00	0.00	0.00	289
farfalla	0.08	0.25	0.12	422
gallina	0.11	0.01	0.03	619
gatto	0.00	0.00	0.00	333
mucca	0.00	0.00	0.00	373
pecora	0.00	0.00	0.00	364
ragno	0.18	0.70	0.29	964
scoiattolo	0.00	0.00	0.00	372
accuracy			0.15	5232
macro avg	0.05	0.10	0.05	5232
weighted avg	0.08	0.15	0.07	5232

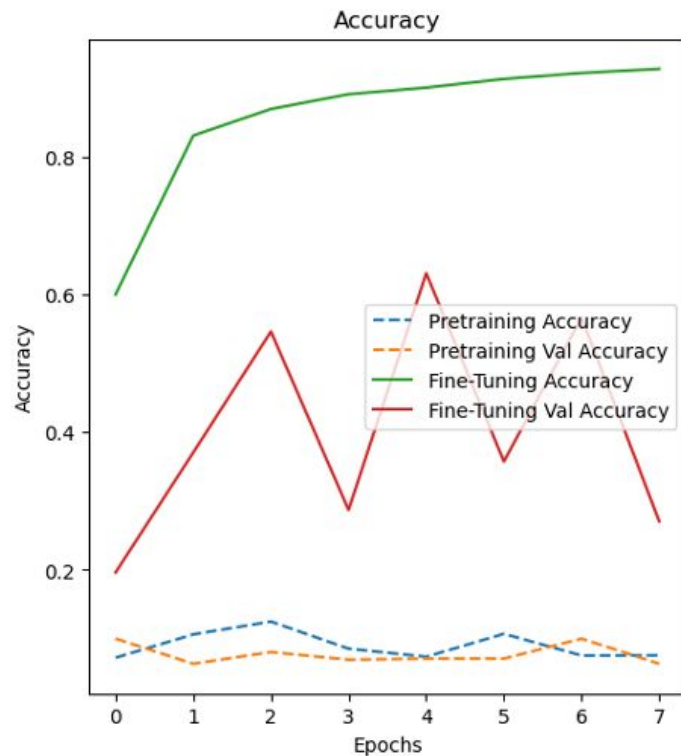
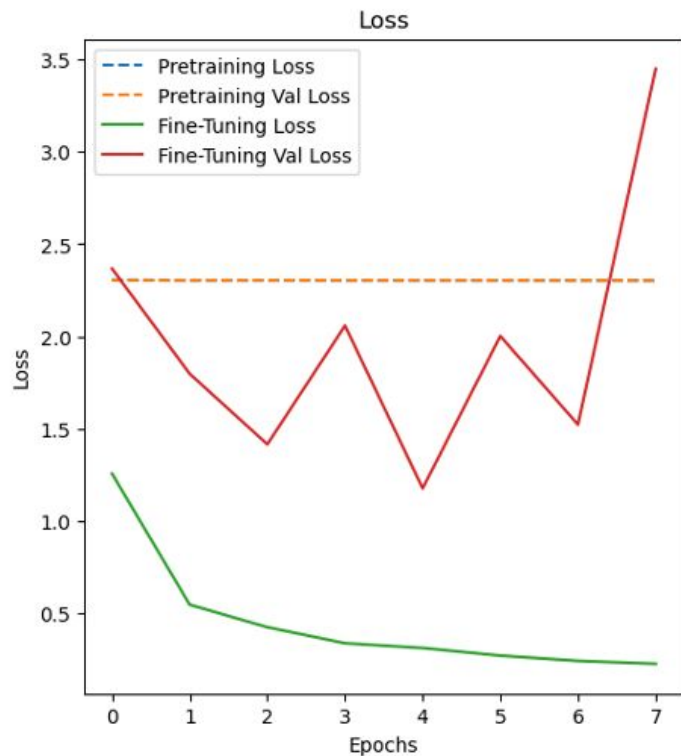
Confusion Matrix:

```
[[ 14   3   2 239  16   2   2   5 685   4]
 [   9   1   2 126   8   2   0   0 375   1]
 [   4   0   0  65   4   0   0   0 213   3]
 [   7   3   0 104  16   0   0   1 291   0]
 [  15   3   1 139   9   1   2   1 447   1]
 [   8   2   0  86   4   0   0   0 233   0]
 [  11   1   0  78   5   3   0   1 273   1]
 [   7   3   0  87   1   3   0   0 263   0]
 [  27   6   2 235  15   1   1   1 675   1]
 [  14   0   0 101   4   2   0   0 251   0]]
```

Model Deployment



Model Deployment



Model Deployment

